

# Problem Statement

Although different banks set their own approval standards, they generally look out for the same type of customers. From machine learning, we can predict the indicators that will drive a successful credit card application.

## Import necessary libraries

```
In [302... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

## Import Dataset

```
In [303... #data set from: https://www.kaggle.com/datasets/samuelcortinhas/credit-card-approval

df = pd.read_csv('/Users/alisonc/Documents/PERSONAL/Data Science Bootcamp/credit_c
```

## Data cleaning & EDA

```
In [304... #get info

df.shape
df.head()
df.info() #check for strings that might relate to the approval
df.describe()
```

Out[304]: (690, 16)

Out[304]:

	Gender	Age	Debt	Married	BankCustomer	Industry	Ethnicity	YearsEmployed	PriorDefau
0	1	30.83	0.000	1	1	Industrials	White	1.25	
1	0	58.67	4.460	1	1	Materials	Black	3.04	
2	0	24.50	0.500	1	1	Materials	Black	1.50	
3	1	27.83	1.540	1	1	Industrials	White	3.75	
4	1	20.17	5.625	1	1	Industrials	White	1.71	

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Gender                690 non-null    int64
1   Age                   690 non-null    float64
2   Debt                  690 non-null    float64
3   Married               690 non-null    int64
4   BankCustomer          690 non-null    int64
5   Industry              690 non-null    object
6   Ethnicity             690 non-null    object
7   YearsEmployed         690 non-null    float64
8   PriorDefault          690 non-null    int64
9   Employed              690 non-null    int64
10  CreditScore           690 non-null    int64
11  DriversLicense        690 non-null    int64
12  Citizen               690 non-null    object
13  ZipCode               690 non-null    int64
14  Income                690 non-null    int64
15  Approved              690 non-null    int64
dtypes: float64(3), int64(10), object(3)
memory usage: 86.4+ KB
```

Out[304]:

	Gender	Age	Debt	Married	BankCustomer	YearsEmployed	PriorDefau
count	690.000000	690.000000	690.000000	690.000000	690.000000	690.000000	690.000000
mean	0.695652	31.514116	4.758725	0.760870	0.763768	2.223406	0.52318
std	0.460464	11.860245	4.978163	0.426862	0.425074	3.346513	0.49982
min	0.000000	13.750000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	22.670000	1.000000	1.000000	1.000000	0.165000	0.000000
50%	1.000000	28.460000	2.750000	1.000000	1.000000	1.000000	1.000000
75%	1.000000	37.707500	7.207500	1.000000	1.000000	2.625000	1.000000

Legend

Gender: 0=Male 1=Female  
Debt: Outstanding Debt  
Married: 0=Single/Divorced 1=Married  
BankCustomer: 0=Not customer 1=Customer  
PriorDefault: 0=Not employed 1=Employed  
DriversLicense: 0=No license 1=With license  
Approved: 0=Not approved 1=Approved

### Rename column header for better understanding

In [305...

```
df = df.rename({'Debt': 'OutsDebt', 'Approved': 'ApprovalStatus'}, axis=1)
```

### Identify the objects and find out the value count of them

In [306...

```
string_cols = df.select_dtypes('object').columns

string_cols
```

```

value_count = df[string_cols].apply(pd.value_counts, normalize=True)
value_count
value_count.index # preview to multiindex, you also get this when doing df.groupby

# we don't need such high granularity usually, so do in loop
for col in string_cols:
    df[col].value_counts()

```

Out[306]: Index(['Industry', 'Ethnicity', 'Citizen'], dtype='object')

Out[306]:

	Industry	Ethnicity	Citizen
<b>Asian</b>	NaN	0.085507	NaN
<b>Black</b>	NaN	0.200000	NaN
<b>ByBirth</b>	NaN	NaN	0.905797
<b>ByOtherMeans</b>	NaN	NaN	0.082609
<b>CommunicationServices</b>	0.055072	NaN	NaN
<b>ConsumerDiscretionary</b>	0.085507	NaN	NaN
<b>ConsumerStaples</b>	0.078261	NaN	NaN
<b>Education</b>	0.036232	NaN	NaN
<b>Energy</b>	0.211594	NaN	NaN
<b>Financials</b>	0.073913	NaN	NaN
<b>Healthcare</b>	0.076812	NaN	NaN
<b>Industrials</b>	0.092754	NaN	NaN
<b>InformationTechnology</b>	0.059420	NaN	NaN
<b>Latino</b>	NaN	0.082609	NaN
<b>Materials</b>	0.113043	NaN	NaN
<b>Other</b>	NaN	0.040580	NaN
<b>Real Estate</b>	0.043478	NaN	NaN
<b>Research</b>	0.014493	NaN	NaN
<b>Temporary</b>	NaN	NaN	0.011594
<b>Transport</b>	0.004348	NaN	NaN
<b>Utilities</b>	0.055072	NaN	NaN
<b>White</b>	NaN	0.591304	NaN

Out[306]: Index(['Asian', 'Black', 'ByBirth', 'ByOtherMeans', 'CommunicationServices', 'ConsumerDiscretionary', 'ConsumerStaples', 'Education', 'Energy', 'Financials', 'Healthcare', 'Industrials', 'InformationTechnology', 'Latino', 'Materials', 'Other', 'Real Estate', 'Research', 'Temporary', 'Transport', 'Utilities', 'White'], dtype='object')

Out[306]:

Energy	146
Materials	78
Industrials	64
ConsumerDiscretionary	59
ConsumerStaples	54
Healthcare	53
Financials	51
InformationTechnology	41
CommunicationServices	38
Utilities	38
Real Estate	30
Education	25
Research	10
Transport	3

Name: Industry, dtype: int64

Out[306]:

White	408
Black	138
Asian	59
Latino	57
Other	28

Name: Ethnicity, dtype: int64

Out[306]:

ByBirth	625
ByOtherMeans	57
Temporary	8

Name: Citizen, dtype: int64

```
In [307... from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

for col in df:
    if df[col].dtypes=='object':
        df[col]=le.fit_transform(df[col])

df
```

Out[307]:

	Gender	Age	OutsDebt	Married	BankCustomer	Industry	Ethnicity	YearsEmployed	Prior
0	1	30.83	0.000	1	1	7	4	1.25	
1	0	58.67	4.460	1	1	9	1	3.04	
2	0	24.50	0.500	1	1	9	1	1.50	
3	1	27.83	1.540	1	1	7	4	3.75	
4	1	20.17	5.625	1	1	7	4	1.71	
...	...	...	...	...	...	...	...	...	...
685	1	21.08	10.085	0	0	3	1	1.25	
686	0	22.67	0.750	1	1	4	4	2.00	
687	0	25.25	13.500	0	0	6	2	2.00	
688	1	17.92	0.205	1	1	2	4	0.04	
689	1	35.00	3.375	1	1	4	1	8.29	

690 rows × 16 columns



### Check for duplicates

```
In [308... df.duplicated().sum()

#Result: 0 duplicates to drop
```

```
Out[308]: 0
```

Results: No duplicates to drop

## Check for missing data

```
In [309... df.isnull().sum().sum()

#Result: 0 null to fill
```

```
Out[309]: 0
```

Results: No missing data to fill

## Observe the mean of the variants compared to the approval status

```
In [310... df.groupby('ApprovalStatus').mean()
```

```
Out[310]:
```

	Gender	Age	OutsDebt	Married	BankCustomer	Industry	Ethnicity	Year
<b>ApprovalStatus</b>								
0	0.707572	29.773029	3.839948	0.691906	0.691906	4.913838	2.953003	
1	0.680782	33.686221	5.904951	0.846906	0.853420	6.302932	2.726384	

### Basic observation through mean

- The mean of Gender and Age in relations to the approval status does not have a big variance. It does not seem to affect approval status.
- A higher Outstanding Debt did not seem to affect the approval rate. It could be that the bank also consider on time debt repayment, as long as the repayment is done on time, it will not affect the approval rate.
- Being Married or a Customer of the bank have very little affect on the approval rate as well.
- Longer Years Employed, Being Employed, good Credit Score and a high income does seem to have push more approvals
- Having Prior Payment seem to lower the approval rate
- Having a Drivers License and Zipcode location dont seem to have any relations to the approval status

## Use Pairplot to identify correlations of the different variants

```
In [313... sns.pairplot(df, hue = 'ApprovalStatus')
plt.show()
```

```
Out[313]: <seaborn.axisgrid.PairGrid at 0x1a491adc2b0>
```



## Identify correlations of the variants through pairplot

- Gender, Marital Status, Citizen, Bank Customer, Employed, Ethnicity and Drivers License does not seem to have any correlations to any other variants
- Age, Industry, Outstanding Debt, Years Employed, Prior Default, Credit Score, Zipcode and Income seem to have come correeslations to each other that we can explore and understand better.

## Further use of heatmap to observe correlations

In [314...

```
df.corr()
```

Out[314]:

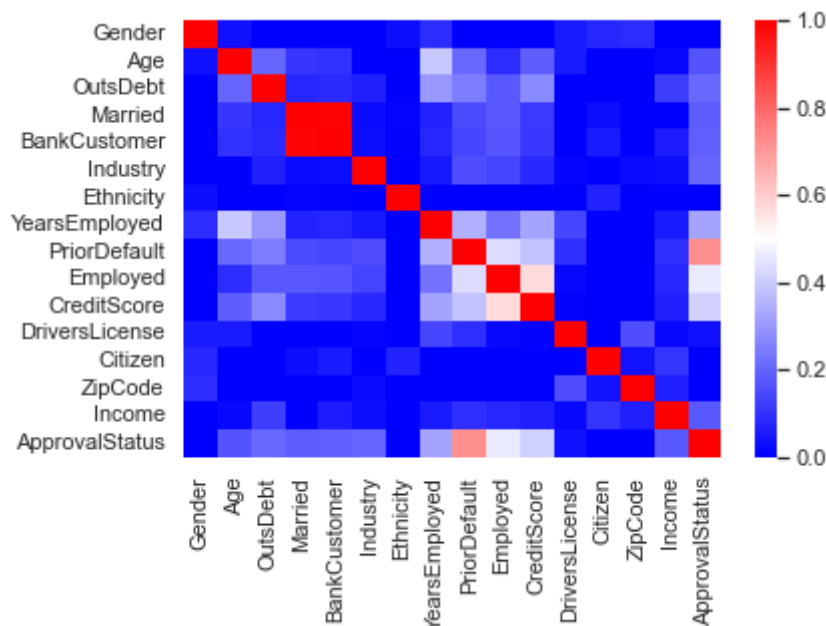
	Gender	Age	OutsDebt	Married	BankCustomer	Industry	Ethnicity	Y
<b>Gender</b>	1.000000	0.035044	-0.041746	-0.068062	-0.071250	-0.111889	0.029492	
<b>Age</b>	0.035044	1.000000	0.202177	0.106929	0.099477	-0.038746	-0.179534	
<b>OutsDebt</b>	-0.041746	0.202177	1.000000	0.074649	0.083781	0.064553	-0.075789	
<b>Married</b>	-0.068062	0.106929	0.074649	1.000000	0.992033	0.021652	0.008226	
<b>BankCustomer</b>	-0.071250	0.099477	0.083781	0.992033	1.000000	0.024677	0.006648	
<b>Industry</b>	-0.111889	-0.038746	0.064553	0.021652	0.024677	1.000000	-0.013881	
<b>Ethnicity</b>	0.029492	-0.179534	-0.075789	0.008226	0.006648	-0.013881	1.000000	
<b>YearsEmployed</b>	0.086544	0.391464	0.298902	0.069945	0.075905	0.048689	-0.177111	
<b>PriorDefault</b>	-0.026047	0.204434	0.244317	0.145073	0.138535	0.154645	-0.114148	
<b>Employed</b>	-0.077784	0.086037	0.174846	0.175428	0.170268	0.134769	-0.030250	
<b>CreditScore</b>	-0.024630	0.187327	0.271207	0.113968	0.111077	0.080107	-0.068072	
<b>DriversLicense</b>	0.051674	0.053599	-0.013023	-0.009784	-0.002402	0.011010	-0.035688	
<b>Citizen</b>	0.075413	-0.006481	-0.116975	0.024319	0.052141	-0.122211	0.070235	
<b>ZipCode</b>	0.086007	-0.078690	-0.217903	-0.017074	-0.009513	0.020551	-0.040003	
<b>Income</b>	-0.002063	0.018719	0.123121	-0.006899	0.057273	0.027820	-0.034251	
<b>ApprovalStatus</b>	-0.028934	0.164086	0.206294	0.180583	0.188964	0.202158	-0.075558	

In [315]...

```
sns.heatmap(df.corr(),
             cmap='bwr',
             vmin=0,
             vmax=1
            )
plt.show()
```

Out[315]:

&lt;AxesSubplot:&gt;



- Looking at both the pair plot as well as the heatmap, we can roughly determine the variables that most affect the approval status. - Namely, Age, Industry, Outstanding Debt, Years Employed, Credit Score and

Income. - We can look at dropping the rest of the columns that does not seem to have a significant effect to the approval rate.

In [316...

df

Out[316]:

	Gender	Age	OutsDebt	Married	BankCustomer	Industry	Ethnicity	YearsEmployed	Prior
0	1	30.83	0.000	1	1	7	4	1.25	
1	0	58.67	4.460	1	1	9	1	3.04	
2	0	24.50	0.500	1	1	9	1	1.50	
3	1	27.83	1.540	1	1	7	4	3.75	
4	1	20.17	5.625	1	1	7	4	1.71	
...	...	...	...	...	...	...	...	...	...
685	1	21.08	10.085	0	0	3	1	1.25	
686	0	22.67	0.750	1	1	4	4	2.00	
687	0	25.25	13.500	0	0	6	2	2.00	
688	1	17.92	0.205	1	1	2	4	0.04	
689	1	35.00	3.375	1	1	4	1	8.29	

690 rows × 16 columns

In [317...

df = df.drop(['Gender', 'BankCustomer', 'Married', 'DriversLicense', 'Citizen', 'Z:

In [318...

df

Out[318]:

	Age	OutsDebt	Industry	YearsEmployed	PriorDefault	Employed	CreditScore	Income	Ap
0	30.83	0.000	7	1.25	1	1	1	0	
1	58.67	4.460	9	3.04	1	1	6	560	
2	24.50	0.500	9	1.50	1	0	0	824	
3	27.83	1.540	7	3.75	1	1	5	3	
4	20.17	5.625	7	1.71	1	0	0	0	
...	...	...	...	...	...	...	...	...	...
685	21.08	10.085	3	1.25	0	0	0	0	
686	22.67	0.750	4	2.00	0	1	2	394	
687	25.25	13.500	6	2.00	0	1	1	1	
688	17.92	0.205	2	0.04	0	0	0	750	
689	35.00	3.375	4	8.29	0	0	0	0	

690 rows × 9 columns

# How does Age affect Approval Status



Pulling out potential correlations between different variants from the pairplot earlier

```
In [319]: fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(x=df['ApprovalStatus'], height=df['Age'], color = 'maroon',
        width = 0.2)

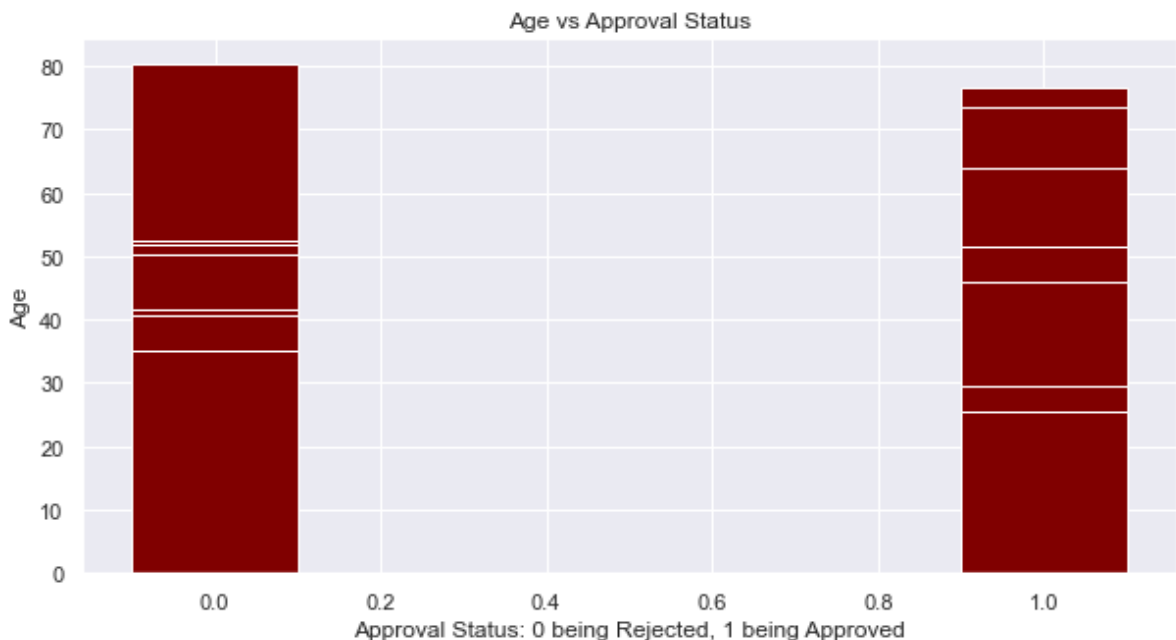
plt.xlabel('Approval Status: 0 being Rejected, 1 being Approved')
plt.ylabel('Age')
plt.title('Age vs Approval Status')
plt.show()
```

Out[319]: <BarContainer object of 690 artists>

Out[319]: Text(0.5, 0, 'Approval Status: 0 being Rejected, 1 being Approved')

Out[319]: Text(0, 0.5, 'Age')

Out[319]: Text(0.5, 1.0, 'Age vs Approval Status')



- From the bar chart, age does not seem to affect approval status

## How does Credit Score affect Approval Status

```
In [320]: fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(x=df['ApprovalStatus'], height=df['CreditScore'], color = 'blue',
        width = 0.2)

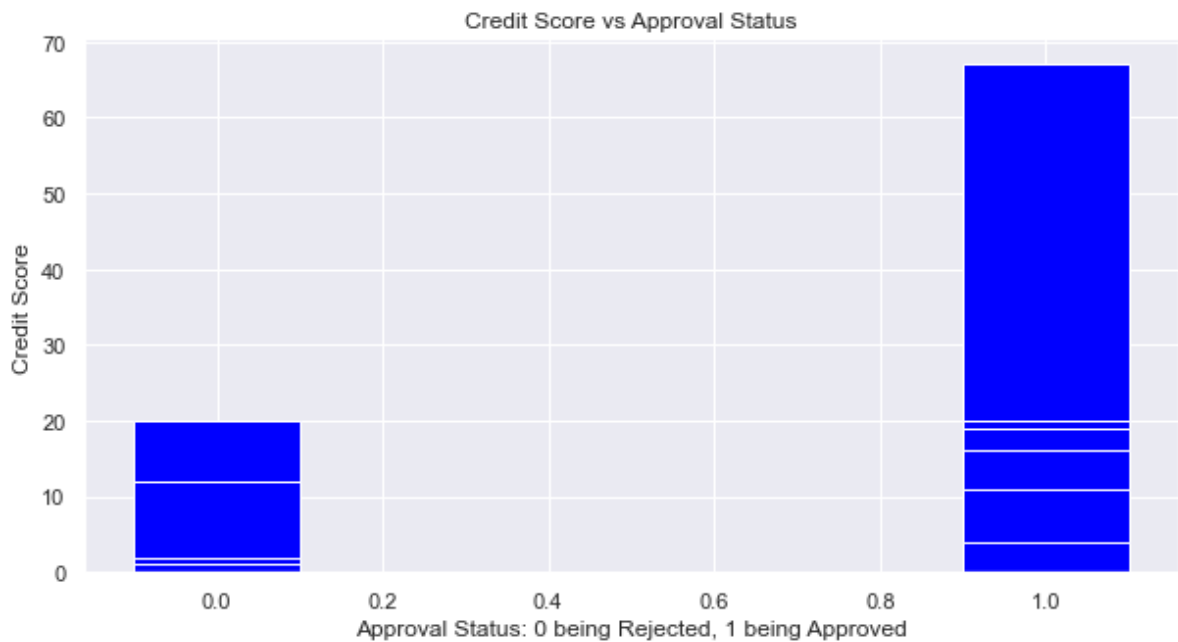
plt.xlabel('Approval Status: 0 being Rejected, 1 being Approved')
plt.ylabel('Credit Score')
plt.title('Credit Score vs Approval Status')
plt.show()
```

Out[320]: <BarContainer object of 690 artists>

Out[320]: Text(0.5, 0, 'Approval Status: 0 being Rejected, 1 being Approved')

Out[320]: Text(0, 0.5, 'Credit Score')

Out[320]: Text(0.5, 1.0, 'Credit Score vs Approval Status')



- We can see that when credit score is higher, there is higher approval rate

## How does Prior Default affect Approval Status

In [321]...

```
sns.set_theme(style='darkgrid')

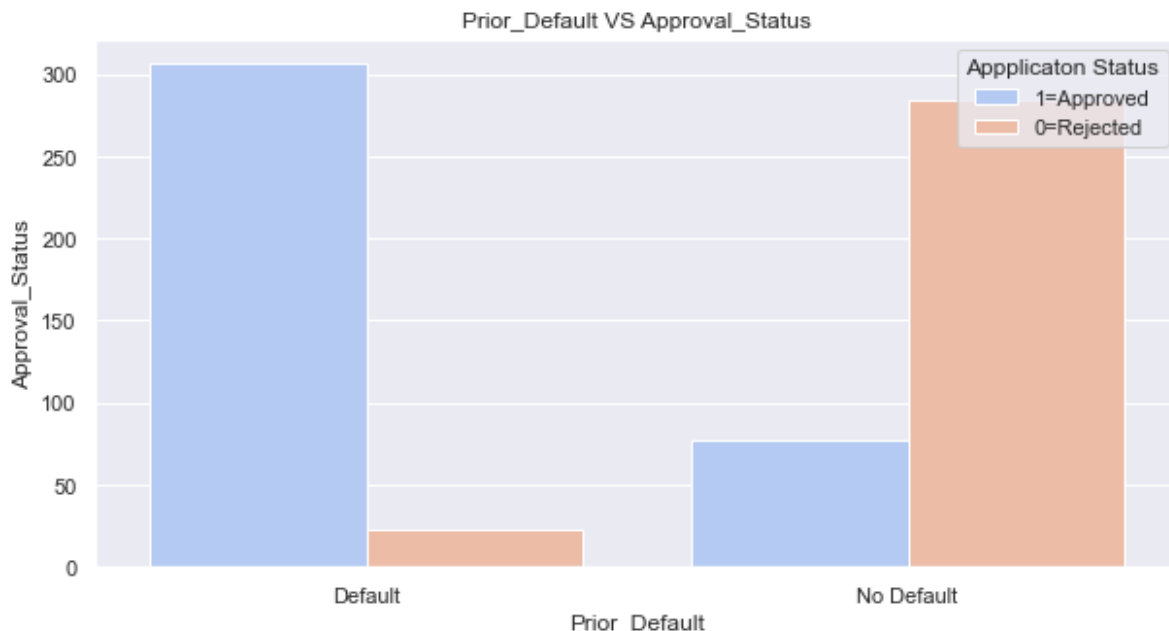
fig = plt.subplots(figsize=(10,5))
PriorDefaultPlot = sns.countplot(data=df, x='PriorDefault', hue='ApprovalStatus', palette='magma')
PriorDefaultPlot.set(xlabel='Prior_Default', ylabel='Approval_Status')
PriorDefaultPlot.set_xticklabels(['Default', 'No Default'])
plt.legend(title='Appplication Status', labels=['1=Approved', '0=Rejected'], loc='upper right')
plt.title('Prior_Default VS Approval_Status')
plt.show()
```

Out[321]: [Text(0.5, 0, 'Prior\_Default'), Text(0, 0.5, 'Approval\_Status')]

Out[321]: [Text(0, 0, 'Default'), Text(1, 0, 'No Default')]

Out[321]: <matplotlib.legend.Legend at 0x1a4b3c11f10>

Out[321]: Text(0.5, 1.0, 'Prior\_Default VS Approval\_Status')



In [ ]: - When there is no default, there is higher approval rate and when there are default.

## How does Income affect Approval Status

```
In [322]: fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(x=df['ApprovalStatus'], height=df['Income'], color = 'green',
        width = 0.2)

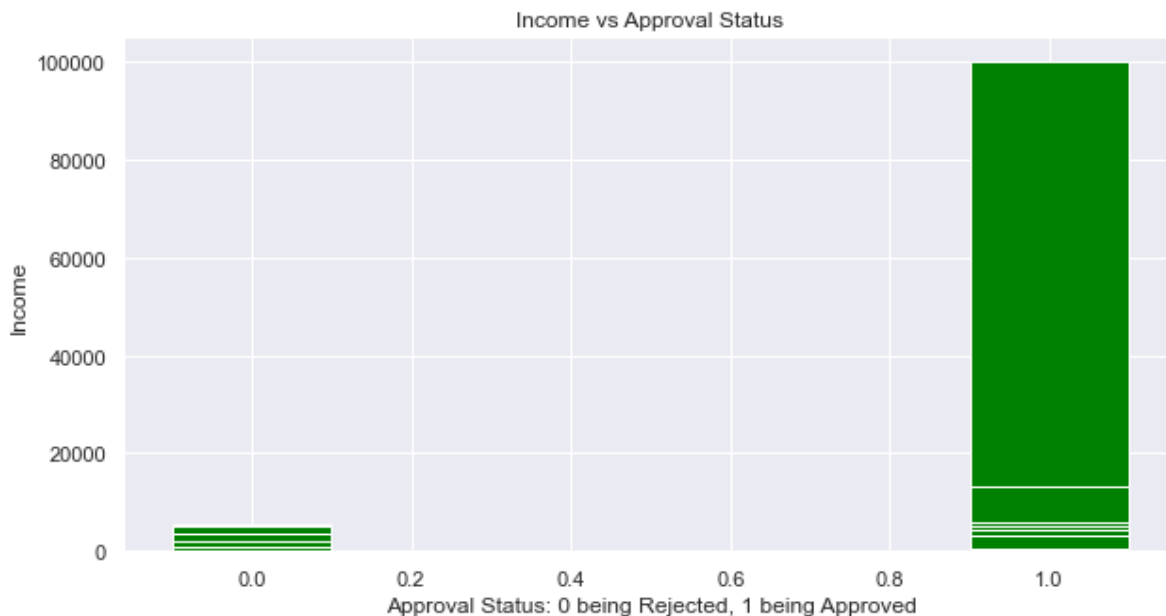
plt.xlabel('Approval Status: 0 being Rejected, 1 being Approved')
plt.ylabel('Income')
plt.title('Income vs Approval Status')
plt.show()
```

Out[322]: <BarContainer object of 690 artists>

Out[322]: Text(0.5, 0, 'Approval Status: 0 being Rejected, 1 being Approved')

Out[322]: Text(0, 0.5, 'Income')

Out[322]: Text(0.5, 1.0, 'Income vs Approval Status')



- High income increases approval rate

## How does being Employed affect Approval Status

```
In [323... sns.set_theme(style='darkgrid')

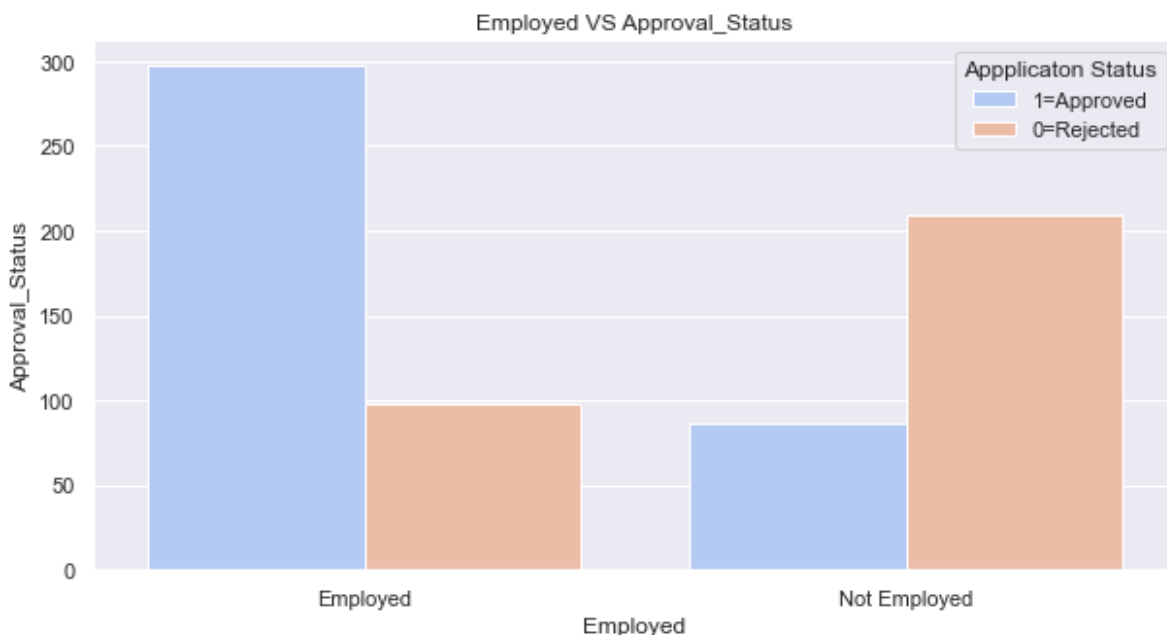
fig = plt.subplots(figsize=(10,5))
EmployedPlot = sns.countplot(data=df,x='Employed', hue='ApprovalStatus', palette='magma')
EmployedPlot.set(xlabel='Employed', ylabel='Approval_Status')
EmployedPlot.set_xticklabels(['Employed', 'Not Employed'])
plt.legend(title='Appplication Status', labels=['1=Approved', '0=Rejected'], loc='upper right')
plt.title('Employed VS Approval_Status')
plt.show()
```

Out[323]: [Text(0.5, 0, 'Employed'), Text(0, 0.5, 'Approval\_Status')]

Out[323]: [Text(0, 0, 'Employed'), Text(1, 0, 'Not Employed')]

Out[323]: <matplotlib.legend.Legend at 0x1a4b1a34fd0>

Out[323]: Text(0.5, 1.0, 'Employed VS Approval\_Status')



- Being employed increases the approval rate compared to not being employed

## How does being Years Employed affect Approval Status

```
In [324]: fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(x=df['ApprovalStatus'], height=df['YearsEmployed'], color = 'yellow',
        width = 0.2)

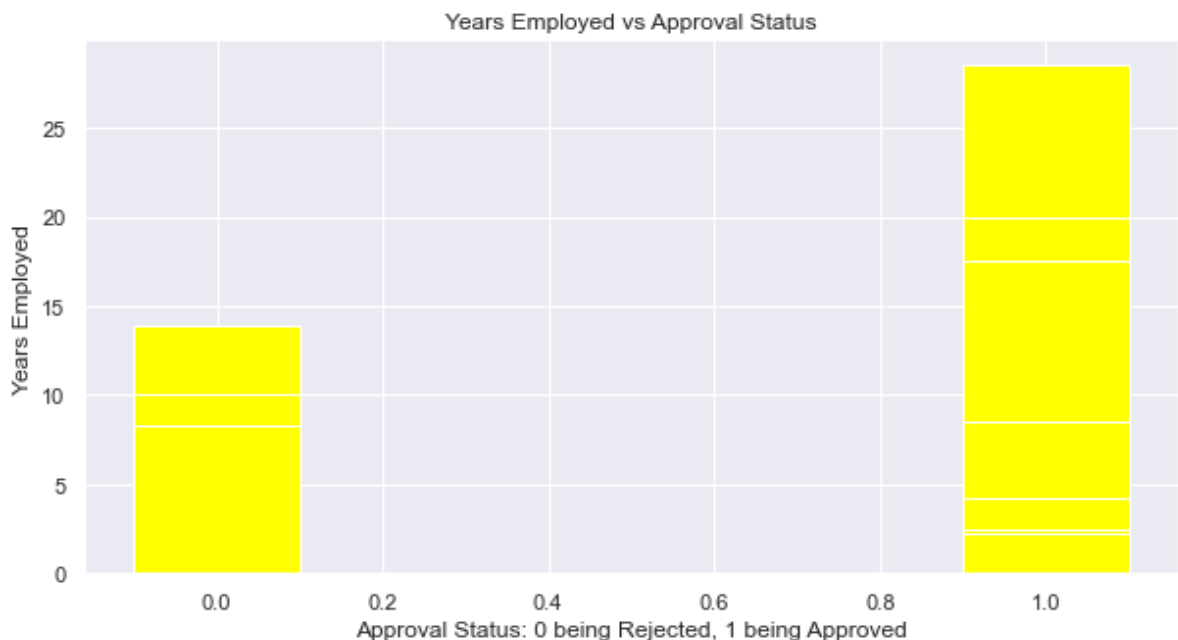
plt.xlabel('Approval Status: 0 being Rejected, 1 being Approved')
plt.ylabel('Years Employed')
plt.title('Years Employed vs Approval Status')
plt.show()
```

Out[324]: <BarContainer object of 690 artists>

Out[324]: Text(0.5, 0, 'Approval Status: 0 being Rejected, 1 being Approved')

Out[324]: Text(0, 0.5, 'Years Employed')

Out[324]: Text(0.5, 1.0, 'Years Employed vs Approval Status')



- Longer history of employment increases approval rate

## How does being Industry affect Approval Status

```
In [292]: fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(x=df['ApprovalStatus'], height=df['Industry'], color = 'pink',
        width = 0.2)

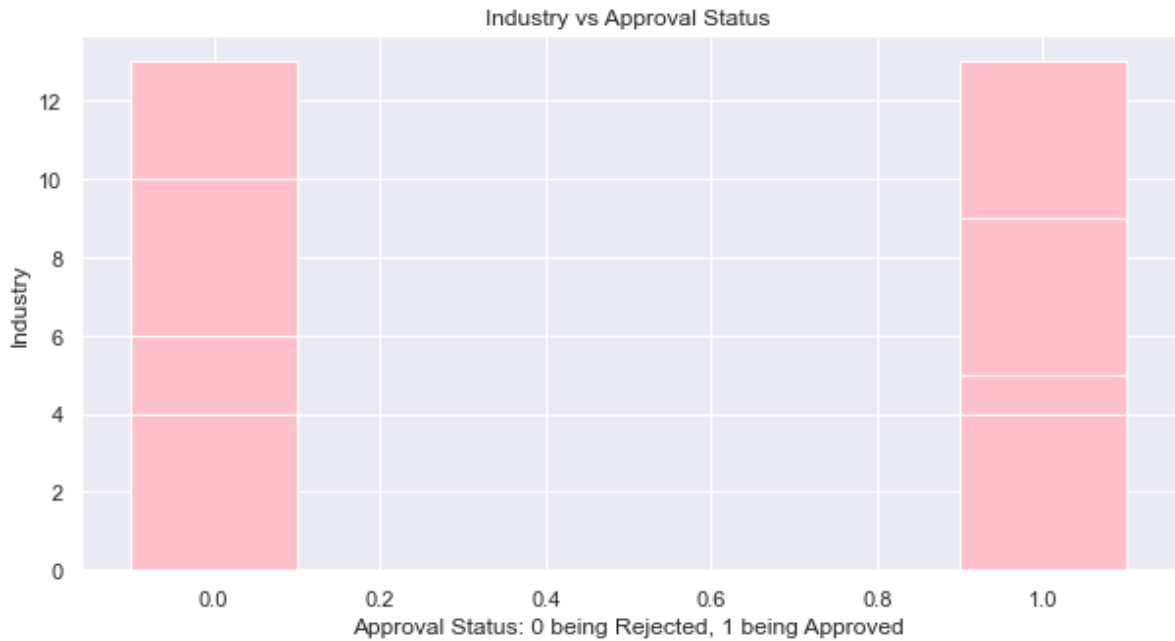
plt.xlabel('Approval Status: 0 being Rejected, 1 being Approved')
plt.ylabel('Industry')
plt.title('Industry vs Approval Status')
plt.show()
```

Out[292]: <BarContainer object of 690 artists>

Out[292]: Text(0.5, 0, 'Approval Status: 0 being Rejected, 1 being Approved')

Out[292]: Text(0, 0.5, 'Industry')

Out[292]: Text(0.5, 1.0, 'Industry vs Approval Status')



- Approval status does not seem to related to industry the person is working in

## How does being Years Outstanding Debt affect Approval Status

```
In [325... fig = plt.figure(figsize = (10, 5))

# creating the bar plot
plt.bar(x=df['ApprovalStatus'], height=df['OutsDebt'], color = 'purple',
        width = 0.2)

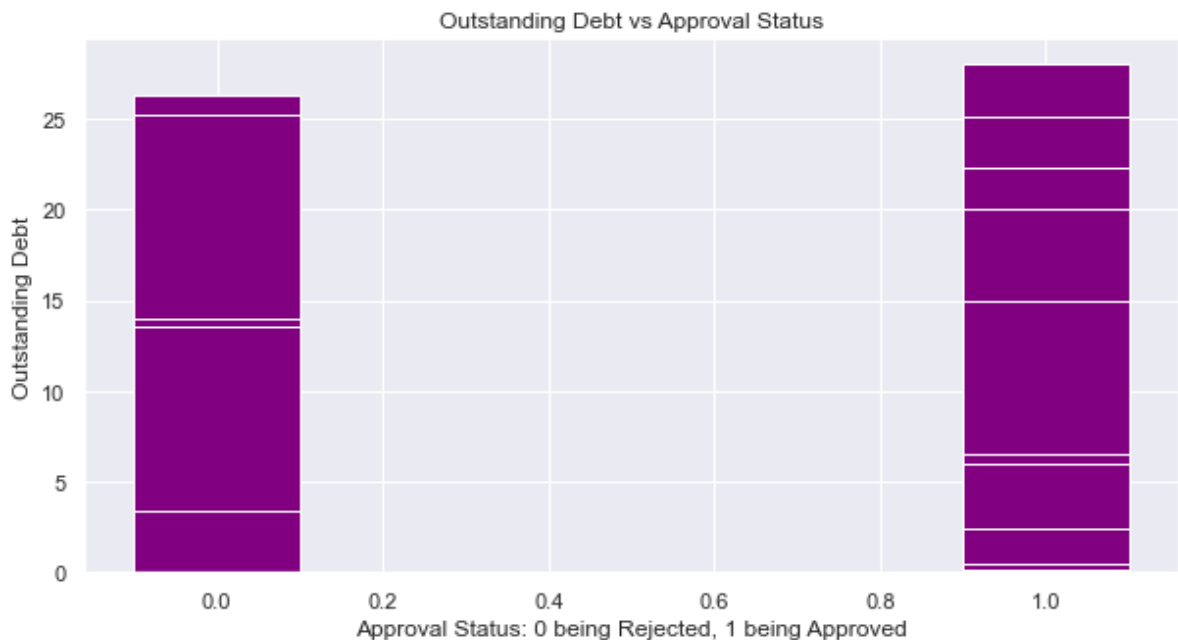
plt.xlabel('Approval Status: 0 being Rejected, 1 being Approved')
plt.ylabel('Outstanding Debt')
plt.title('Outstanding Debt vs Approval Status')
plt.show()
```

Out[325]: <BarContainer object of 690 artists>

Out[325]: Text(0.5, 0, 'Approval Status: 0 being Rejected, 1 being Approved')

Out[325]: Text(0, 0.5, 'Outstanding Debt')

Out[325]: Text(0.5, 1.0, 'Outstanding Debt vs Approval Status')



- Surprised that outstanding debt does not affect approval status much. The bank likely also take into account the timely repayment. If repayment is done on time, the amount of debt does not affect the status.

## Further Clean the Data

Drop column that does not affect the Approval Status

```
In [326...] df = df.drop(['Industry', 'OutsDebt', 'Age'], axis = 1)
```

```
In [327...] df
```

```
Out[327]:
```

	YearsEmployed	PriorDefault	Employed	CreditScore	Income	ApprovalStatus
0	1.25	1	1	1	0	1
1	3.04	1	1	6	560	1
2	1.50	1	0	0	824	1
3	3.75	1	1	5	3	1
4	1.71	1	0	0	0	1
...	...	...	...	...	...	...
685	1.25	0	0	0	0	0
686	2.00	0	1	2	394	0
687	2.00	0	1	1	1	0
688	0.04	0	0	0	750	0
689	8.29	0	0	0	0	0

690 rows × 6 columns

## Machine Learning

Determine the classifier with the highest accuracy

## Knearest neighbors

```
In [328... #import libraries for KNN
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

```
In [329... # y=df.ApprovalStatus
# X=df.drop('ApprovalStatus',axis=1)
```

```
In [334... # STEP 1: split X and y into training and testing sets (using random_state for reproducibility)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=99)

# STEP 2: train the model on the training set (using K=1)
knn = KNeighborsClassifier(n_neighbors=1)

knn.fit(X_train, y_train) #first stage is always train, dont use test

# STEP 3: test the model on the testing set, and check the accuracy
y_pred_class = knn.predict(X_test)
print(accuracy_score(y_test, y_pred_class))
```

```
Out[334]: KNeighborsClassifier(n_neighbors=1)
0.5606936416184971
```

## Optimise hyperparameter

```
In [335... # test with 50 neighbors
knn = KNeighborsClassifier(n_neighbors=50)
knn.fit(X_train, y_train)
y_pred_class = knn.predict(X_test)
print(accuracy_score(y_test, y_pred_class))
```

```
Out[335]: KNeighborsClassifier(n_neighbors=50)
0.6705202312138728
```

```
In [338... # test with 80 neighbors
knn = KNeighborsClassifier(n_neighbors=80)
knn.fit(X_train, y_train)
y_pred_class = knn.predict(X_test)
print(accuracy_score(y_test, y_pred_class))
```

```
Out[338]: KNeighborsClassifier(n_neighbors=80)
0.6589595375722543
```

```
In [342... # test with 46 neighbors
knn = KNeighborsClassifier(n_neighbors=46)
knn.fit(X_train, y_train)
y_pred_class = knn.predict(X_test)
print(accuracy_score(y_test, y_pred_class))
```

```
Out[342]: KNeighborsClassifier(n_neighbors=46)
0.653179190751445
```

Highest Accuracy using n\_neighbors=50: ~65.31%



# Logistic Regression

```
In [343... df_standardized = (df-df.mean())/df.std()
df_standardized.mean()
df_standardized.std()
```

```
Out[343]: YearsEmployed      1.879334e-16
PriorDefault      -5.399867e-16
Employed          -7.852012e-16
CreditScore       2.067589e-16
Income           -1.448117e-18
ApprovalStatus    -5.766724e-16
dtype: float64
```

```
Out[343]: YearsEmployed      1.0
PriorDefault      1.0
Employed          1.0
CreditScore       1.0
Income           1.0
ApprovalStatus    1.0
dtype: float64
```

```
In [348... from sklearn.linear_model import LogisticRegression

from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score

logreg = LogisticRegression()

X = df.drop('ApprovalStatus',axis=1)

y = df.ApprovalStatus

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

logreg.fit(X_train, y_train)

y_pred = logreg.predict(X_test)

print('Accuracy: ', accuracy_score(y_test, y_pred))
```

C:\Users\alisonc\Documents\DataScience\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[348]: LogisticRegression()
```

```
Accuracy: 0.8728323699421965
```

Accuracy: ~87.28%

## Decision Tree

```
In [349... X = df.drop('ApprovalStatus',axis=1)

y = df.ApprovalStatus
```

```
from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier()

#still using the same train size 80%, test size 20%
X_train, X_test, y_train, y_test = train_test_split(X,y, random_state=2, test_size=0.2)

#scaling the features
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_X_train = scaler.fit_transform(X_train)
scaled_X_test = scaler.fit_transform(X_test)

dt.fit(scaled_X_train,y_train)

y_preds = dt.predict(scaled_X_test)

accuracy_score(y_test,y_preds)
```

Out[349]: DecisionTreeClassifier()

Out[349]: 0.8043478260869565

Accuracy: ~80.43%

## Summary

- KNN, Logistic Regression and Decision Tree were applied to predict a new customer's credit card approval rate. - Among all 3 classifiers, Logistic Regression's accuracy is the highest. - I feel like i might have dropped too many columns while doing EDA. I should consider exploring different chart types to plot different types of variants to make a better decision.